

Machine Learning with Go

Ephraim Berkovitch, January 2021



Introduction

- + Ephraim Berkovitch - senior software engineer in ZipRecruiter
 - Building full-stack solutions in Go and Python
 - Contributes to open-source (HaSadna, Ben-Yehuda Project, Yiddish apps)
- + By the end of this session you will understand
 - how simple ML algorithms work
 - how to run them in Go
 - how different ML Golang libraries compare with their Python counterparts

Agenda

- + Why Go?
- + ML using Go vs Python
- + ML Steps
- + Gorgonia
- + Flower species prediction by classification
- + Linear regression
- + Gota - counterpart of pandas
- + Gonum - counterpart of numpy and matplotlib
- + Gophernotes - Jupyter kernel for Go

Accompanying Materials

- + https://github.com/ephraimberkovitch/golang_ml
- +
- + <https://github.com/go-gota/gota>
- + <https://github.com/gorgonia/gorgonia>
- + <https://github.com/gonum/gonum>

Why Go?

- + simple syntax
 - o clearly describe complex algorithms
 - o does not obscure developers from understand how to run efficient optimized code
- + easy to deploy and code
- + easy to understand and debug
- + have its performance measured
- + great tooling
 - o test coverage
 - o presentations

ML using Go vs Python


- + ML frameworks use languages like C/C++ for actual computation
 - o all offer a Python interface
- + To actually run a production machine learning API at scale, we need an infrastructure to
 - o Autoscaling, so that traffic fluctuations don't break our APIs (and our AWS stays manageable)
 - o API management, to handle multiple deployments
 - o Rolling updates, so that we can update models while still serving requests
- + Reliability - first line of defense, to avoid silly type errors - static typing, compilation step, no dead code, no unused imports
- + **Python for data science and scripting, Go for infrastructure**

ML Steps

- + Gather data
- + Perform exploratory data analysis
- + Determine the correct machine learning solution to use
- + Build a model
- + Train the model
- + Test the model

Gorgonia

- + Deep learning in Go
 - go get gorgonia.org/gorgonia
- + Write and evaluate math equations with multi dimensional arrays
 - `g := gorgonia.NewGraph() // expression/computation graph`
- + VM - understands graphs and computes it
 - compile-once, run-many-times
- + *Demonstration 1 - Gorgonia 1st steps*



Flower Species Prediction by Classification


+ Iris dataset (1935) - <https://gist.github.com/curran/a08a1080b88344b0c8a7>

+ features - X

- sepal_length - אורך עלי הגביע
- sepal_width - רוחב עלי הגביע
- petal_length - אורך עלי הכותרת
- petal_width - רוחב עלי הכותרת

+ target - y

- species - setosa, versicolor, virginica



Flower Species Prediction by Classification



Linear Regression

- + encode the data (the true facts from observation of different flowers) into a matrix
 - o containing 5 columns (sepal length, sepal width, petal length, petal width and 1 for the bias).
- + A row of the matrix represents a flower

$$y = \theta_0 + \theta_1 * \text{sepal_length} + \theta_2 * \text{sepal_width} + \theta_3 * \text{petal_length} + \theta_4 * \text{petal_width}$$

$$x = [\text{sepal_length} \quad \text{sepal_width} \quad \text{petal_length} \quad \text{petal_width} \quad 1]$$

$$\Theta = [\theta_4 \theta_3 \theta_2 \theta_1 \theta_0]$$

$$y = x \cdot \Theta$$

Linear Regression

- + Encode categorical data

- setosa = 1.0
- virginica = 2.0
- versicolor = 3.0

- + Other encoding techniques:

- species_setosa = 1/0: species_virginica=0/1: species_versicolor=0/1

- + Learning phase cost: $cost = \frac{1}{m} \sum_{i=1}^m (X^{(i)} \cdot \Theta - Y^{(i)})^2$

- + Use **gradient descent** to lower the cost

Gota - counterpart of Pandas

- + Data wrangling library
- + Generate / load the training / test set with Gota
 - o dataframe
 - o series
- + Pre-processing
- + Generate training / validation data split
- + *Demonstration 2 - gota 1st steps*

Gonum - counterpart of numpy and matplotlib

- + Scientific and plotting library
 - `go get gonum.org/v1/gonum`
 - `go get gonum.org/v1/plot`
- + *Demonstration 3 - gonum/plot 1st steps*

Demonstrations 4 & 5

+ iris

- training_iris
- cli_iris

+ congress votes (~1980s)

- training_congress
- cli_congress

Jupyter kernel for Go

- + Gophernotes

- `docker run -it -p 8888:8888 -v $(pwd):/usr/share/notebooks gopherdata/gophernotes:latest-ds`

- + *Demonstration 6 - gophernotes*

Comparisons

+ Gorgonia

- <https://golangdocs.com/golang-machine-learning-libraries>

+ Gota

- performance - <https://github.com/go-gota/gota/issues/16>
- other libraries - <https://mungingdata.com/go/dataframes-gota-qframe>

+ gophernotes

- <https://github.com/gopherdata/gophernotes> - backend - interpreter
- alternatives - lgo <https://github.com/yunabe/lgo> - backend - official compiler, elder versions

More

- + TensorFlow for Go is not mature yet
 - o https://www.tensorflow.org/install/lang_go - The TensorFlow Go API is not covered by the TensorFlow API stability guarantees
 - o **galeone/tfgo** has its fixes to official TensorFlow Go bindings
- + Cortex - machine learning model serving infrastructure
 - o <https://github.com/cortexlabs/cortex>
 - o take a trained model and automatically implement all the infra features needed—like reproducible deployments, scalable request handling, automated monitoring, etc—to deploy it as an API

References

- + **Michael Bironneau, Toby Coleman**, *Machine Learning with Go Quick Start Guide*, Packt Publishing, 2019
- + **Xuanyi Chew**, *Go Machine Learning Projects*, Packtpub, 2018
- + <https://towardsdatascience.com/why-we-deploy-machine-learning-models-with-go-not-python-a4e35ec16deb>
- + <https://towardsdatascience.com/understanding-regression-using-covid-19-dataset-detailed-analysis-be7e319e3a50>
- + <https://towardsdatascience.com/go-for-data-science-lets-try-46850b12a189>

References

- + <https://tech.travelaudience.com/training-tensorflow-models-in-python-and-serving-with-go-1b2a9386b0ff>
- + <https://jinglescode.github.io/2019/05/07/why-linear-regression-is-not-suitable-for-classification/>
- + <https://mungingdata.com/go/dataframes-gota-qframe/#:~:text=Go%20has%20great%20DataFrame%20libraries,manipulation%20and%20grouping%20functionality%20natively>
- + <https://towardsdatascience.com/a-gentle-journey-from-linear-regression-to-neural-networks-68881590760e>
- + <https://github.com/RedHatOfficial/GoCourse>